

# Homework 5 – Code & Graphs

Gilberto Feliu

Student ID: 801257813

<https://github.com/gfeliu-rgb/Intro-to-ML-ECGR-4105>

## Problem 1A

Linear Model Code:

```
1 def model(t_u, w, b):
2     return w * t_u + b
```

Updated Code for Problem 1A:

```
1 def model(t_u, w2, w1, b):
2     return w2 * t_u**2 + w1 * t_u + b

1 params = torch.tensor([1.0, 1.0, 0.0], requires_grad=True)
2 learning_rate = 1e-2
3 optimizer = optim.SGD([params], lr=learning_rate)
4
5 training_loop(
6     n_epochs = 5000,
7     optimizer = optimizer,
8     params = params,
9     t_u = t_un,
10    t_c = t_c
11 )
```

## Problem 1B

Changed Code:

```
1 params_init = [1.0, 1.0, 0.0]
2 learning_rates = [0.1, 0.01, 0.001, 0.0001]
3
4 for lr in learning_rates:
5     print("Training with learning rate =", lr)
6     params = torch.tensor(params_init, requires_grad=True)
7     optimizer = optim.SGD([params], lr=lr)
8     trained_params = training_loop(
```

```

9      n_epochs = 5000,
10     optimizer = optimizer,
11     params = params,
12     t_u = t_un,
13     t_c = t_c
14   )
15   print("Final params:", trained_params)
16   print("")

```

## Results:

```

Training with learning rate = 0.1 ... Final loss: nan
Training with learning rate = 0.01 ... Final loss: nan
Training with learning rate = 0.001 ... Final loss: nan
Training with learning rate = 0.0001 ... Final loss: 3.8615105152130127
Final params: tensor([-0.8881, 0.5570, -0.8753])

```

Best nonlinear model learning rate: 0.0001  
 Best nonlinear final loss: 3.8615

## Problem 1C

### Comparison Code:

```

1 def training_loop(n_epochs, optimizer, params, t_u, t_c, model_fn):
2     for epoch in range(1, n_epochs + 1):
3         t_p = model_fn(t_u, *params)
4         loss = loss_fn(t_p, t_c)
5         optimizer.zero_grad()
6         loss.backward()
7         optimizer.step()
8         if epoch % 500 == 0:
9             print('Epoch %d, Loss %f' % (epoch, float(loss)))
10    return params

```

```

1 params_linear = torch.tensor([1.0, 0.0], requires_grad=True)
2 optimizer_linear = optim.SGD([params_linear], lr=0.01)
3 trained_linear_params = training_loop(
4     n_epochs=5000,
5     optimizer=optimizer_linear,
6     params=params_linear,
7     t_u=t_un,
8     t_c=t_c,
9     model_fn=linear_model
10 )
11 pred_linear = linear_model(t_un, *trained_linear_params)
12 pred_nonlinear = model(t_un, *best_params)
13 plt.figure(figsize=(8,5))

```

```

14 plt.scatter(t_un, t_c, color='black', label='Actual data')
15 plt.plot(t_un, pred_linear.detach(), color='blue', label='Linear model')
16 plt.plot(t_un, pred_nonlinear.detach(), color='red', label='Non-linear
    model')
17 plt.legend(); plt.show()

```

**Results:**

```

Linear model final loss: 2.9276480674743652
Best nonlinear model final loss: 3.8615105152130127

```

## Problem 2A

**Source Code:**

```

1 import numpy as np, pandas as pd, torch, torch.optim as optim
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import MinMaxScaler
4
5 file_path = "Housing.csv"
6 df = pd.read_csv(file_path)
7 inputs = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking']
8 X = df[inputs].values
9 Y = df['price'].values
10 X_0 = np.ones((len(Y), 1))
11 X = np.hstack((X_0, X))
12 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2)
13 scaler = MinMaxScaler()
14 X_train[:,1:] = scaler.fit_transform(X_train[:,1:])
15 X_test[:,1:] = scaler.transform(X_test[:,1:])

```

**Results (2A):**

```

Epoch 500 Loss= 1.70e12
Epoch 2000 Loss= 1.44e12
Validation loss: 2.44e12
Trained parameters [B, W1..W5]:
[2508056.5 2933135.8 1395469.2 2351054.8 1603754.9 1423087.1]

```

## Problem 2B

**Results:**

```

Best learning rate = 0.1
Lowest validation loss = 2.29e12 at Epoch 2000
Higher learning rates diverged, lower rates trained slower.

```

## Problem 3A

NN Model:

```
1 class NN(torch.nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.fc1 = torch.nn.Linear(6, 8)
5         self.fc2 = torch.nn.Linear(8, 1)
6     def forward(self, x):
7         x = torch.relu(self.fc1(x))
8         return self.fc2(x)
```

Results (3A):

Training time = 0.333 s  
Final Training Loss = 3.98e12  
Evaluation Accuracy ( $R^2$ ) = -0.3001

## Problem 3B

Extended NN Model:

```
1 class NN(torch.nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.fc1 = torch.nn.Linear(6, 8)
5         self.fc2 = torch.nn.Linear(8, 8)
6         self.fc3 = torch.nn.Linear(8, 8)
7         self.fc4 = torch.nn.Linear(8, 1)
8     def forward(self, x):
9         x = torch.relu(self.fc1(x))
10        x = torch.relu(self.fc2(x))
11        x = torch.relu(self.fc3(x))
12        return self.fc4(x)
```

Results (3B):

Learning rate 0.01 -> Loss = NaN  
Learning rate 1e-9 -> Loss = 2.52e13,  $R^2$  = -4.96  
Training time = 0.277 s

## Graphs

